

## What is a Database?

A database is an organized collection of data that can be easily accessed, managed, and updated. It is used to store information like customer details, sales records, employee data, product catalogs, etc.

Databases are usually managed by Database Management Systems (DBMS), which provide tools to store, retrieve, and manipulate the data.

## Types of Databases

Databases can be classified based on how they store, organize, and access data.

### 1. Relational Databases (RDBMS)

- Store data in tables (rows and columns)
- Use Structured Query Language (SQL) to manage data
- Examples:
  - Oracle Database
  - Microsoft SQL Server
  - MySQL
  - PostgreSQL

### 2. NoSQL Databases

- Used for unstructured or semi-structured data
- Don't use fixed table schemas
- Types of NoSQL databases:
  - Document-based (e.g., MongoDB)
  - Key-Value (e.g., Redis)
  - Column-family (e.g., Cassandra)
  - Graph (e.g., Neo4j)

### 3. In-Memory Databases

- Store data in RAM for ultra-fast access
- Suitable for real-time processing
- Examples: Redis, Memcached

### 4. Cloud Databases

- Hosted on cloud platforms like Azure, AWS, or GCP
- Highly scalable and managed (e.g., Azure SQL, Amazon RDS)

## 5. Object-Oriented Databases

- Store data in the form of objects (like in object-oriented programming)
- Support features like inheritance, polymorphism
- Example: db4o

## 6. Distributed Databases

- Data is stored across multiple locations or servers
- Can be replicated or partitioned
- Improves availability and scalability

## 7. Graph Databases

- Store data in nodes and edges to represent relationships
- Excellent for social networks, recommendation engines
- Example: Neo4j

## Introduction to Neo4j

Neo4j is a **graph database** that is designed to store, query, and manage highly connected data efficiently. Unlike traditional relational databases, Neo4j uses **nodes** and **relationships** to represent and store data.

### Key Features of Neo4j

- **Graph-based Storage:** Stores data as nodes (entities) and relationships (connections) rather than tables.
- **Cypher Query Language (CQL):** A powerful and easy-to-read query language similar to SQL but optimized for graphs.
- **Schema-free Model:** No fixed schema required, allowing flexible data representation.
- **High Performance for Connected Data:** Optimized for traversing relationships efficiently.
- **ACID Compliance:** Ensures reliability, consistency, and integrity of transactions.
- **Built-in Indexing:** Uses native indexes for faster queries.
- **Scalability:** Supports horizontal and vertical scaling.

### Components of Neo4j

1. **Nodes:** Represent entities (e.g., Person, Product, City).
2. **Relationships:** Define connections between nodes (e.g., FRIENDS\_WITH, PURCHASED).
3. **Properties:** Key-value pairs stored in nodes and relationships.

4. **Labels:** Categorize nodes into types (e.g., User, Product).
5. **Indexes & Constraints:** Improve search performance and enforce data rules.

## Basic Cypher Queries

### Creating Nodes

```
CREATE (p:Person {name: 'Alice', age: 30})
```

### Creating Relationships

```
MATCH (a:Person {name: 'Alice'}), (b:Person {name: 'Bob'})
```

```
CREATE (a)-[:FRIENDS_WITH]->(b)
```

### Retrieving Data

```
MATCH (p:Person) RETURN p
```

### Updating Data

```
MATCH (p:Person {name: 'Alice'})
```

```
SET p.age = 31
```

```
RETURN p
```

### Deleting Nodes

```
MATCH (p:Person {name: 'Alice'})
```

```
DETACH DELETE p
```

## When to Use Neo4j?

- **Social Networks:** Managing relationships between users.
- **Recommendation Systems:** Suggesting products, friends, or content.
- **Fraud Detection:** Identifying suspicious connections.
- **Network & IT Management:** Analyzing infrastructure dependencies.
- **Knowledge Graphs:** Storing and linking large datasets.


## Getting Started with Neo4j

1. **Install Neo4j:** Download from [neo4j.com](https://neo4j.com)
2. **Use Neo4j Browser:** Interactive web-based UI to execute Cypher queries.
3. **Try Neo4j Sandbox:** Cloud-hosted environment to experiment with.
4. **Explore Neo4j Aura:** Fully managed cloud database for production use.

**Conclusion** - Neo4j is a powerful tool for handling complex and connected data structures. With its graph-based approach, intuitive Cypher language, and scalability, it is widely used across various industries for analytical and operational applications.


## AWS Cloud

- Global Infra:
  - Regions
  - Availability Zones
  - Global Network



## Azure Cloud

- Global Infra:
  - Regions
  - Zones
  - Backbone Network



- [Refer Here](#) for exploring datacenters of azure

## Cloud Services

- Service: This is what is offered by Cloud Providers. Each service is metered and billing is generally pay as you go.
- Categories:
  - Compute
  - Storage
  - Network
  - Databases
  - AI/ML
  - DevOps
  - Big Data Services
- Each Service will have
  - billing
  - different options
  - Automation options
- On a Generic Note every cloud provides the following approaches to create infra
  - Website
    - AWS: Console
    - Azure: Portal
  - Command Line
    - AWS CLI
    - Azure CLI & Azure Powershell
  - Client Libraries/SDK
    - AWS SDK:
      - python: boto3
      - AWS SDK for Java ...
    - Azure SDK
  - Template
    - AWS Cloudformation
    - Azure ARM Template (Bicep)
    - Terraform (DevOps)
- Personal Accounts : We will be using free tier accounts for learning cloud.
  - By default you have admin access.
- Enterprise Accounts:
  - Landing Zones:
    - Governance rules are established
    - Roles established
    - Generally we will have multiple accounts
    - Templates
  - User Accounts
- Roles:
  - Solutions Architect
  - Developer
  - Cloud Engineer (Ops Engineer)

## Activities around Databases

- Data Modelling

- Capacity Planning
- BCDR (Business Continuity and Disaster Recovery)
  - Backups
  - Recoveries
- Patching:
  - OS
  - DBMS Updates
- Storage needs
  - Expanding sizes
  - Archiving
- Migration

## Databases offered by AWS

- RDS (Relational Database Services):
  - Engines:
    - mysql
    - PostgreSQL
    - Microsoft SQL Server
    - Oracle
    - IBM DB2
    - Aurora
      - mysql
      - postgresql
- DynamoDB (NOSQL)
- Document DB (Mongo DB)
- KeySpaces (Cassandra)
- QLDB (Quantum Ledger Database)
- Timestream:
  - IoT Devices
- Neptune (Graph Databases)
- Elastic Cache (Cache Databases)
- Memory DB (Cache Database)
- NOSQL Databases can scale horizontally using many techniques

- Shards and replicas



## Databases on Azure

- Azure SQL (Microsoft SQL Server)
  - Ledger functionality
- Azure SQL for Postgres
- Azure SQL for mysql
- Azure Cosmos DB
  - Document
  - Key Value
  - Column Databases
  - Graph
- Azure Cache

## Database as a Service

- Database as a service is offering from cloud providers where we get
  - Database directly (rather than installing)
  - Options for
    - automated backups
    - replications
    - patching
    - failovers
- Drawbacks:
  - not all versions are supported



- some features might not be supported.
- Advantages:
  - Minimum Administration
  - Easy configurations
- DBA Responsibilities (On-prem):
  - Backups and Recoveries
  - Replications
  - Patching
  - Performance Tuning
- With cloud databases or Database as a Service there is only one option partially left for DBA's i.e. Performance Tuning
- refer primary index, secondary index, full table scans.

## AWS Relational Database Services (RDS)

- RDS is Database as a service offered by AWS supporting following database engines
  - Microsoft SQL Server
  - mySQL
  - PostgreSQL
  - Oracle
  - IBM DB2

Free Tier 12 MONTHS FREE

Amazon RDS

**750 Hours**


per month of database usage  
(applicable DB engines)

Managed Relational Database  
Service for MySQL,  
PostgreSQL, MariaDB, or SQL  
Server.

**750 Hours** of Amazon RDS Single-  
AZ db.t2.micro, db.t3.micro, and  
db.t4g.micro Instances usage  
running MySQL, MariaDB,  
PostgreSQL databases each month  
(applicable DB engines)

- Free plan: (db.t2.micro/db.t3.micro)
- RDS instances require
  - vpc
  - subnet group

◦ security group



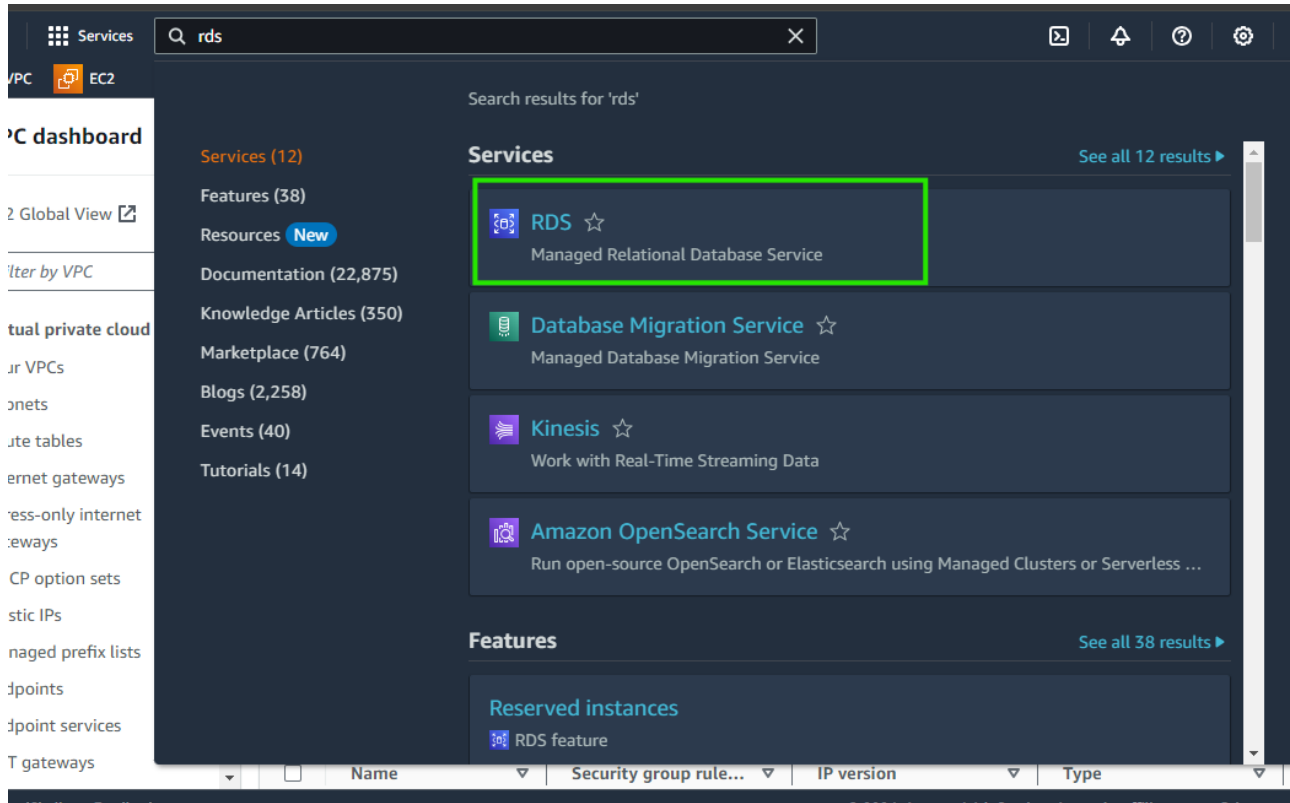
- RDS instances are sized by db instance types

Create a free mysql database in mumbai region

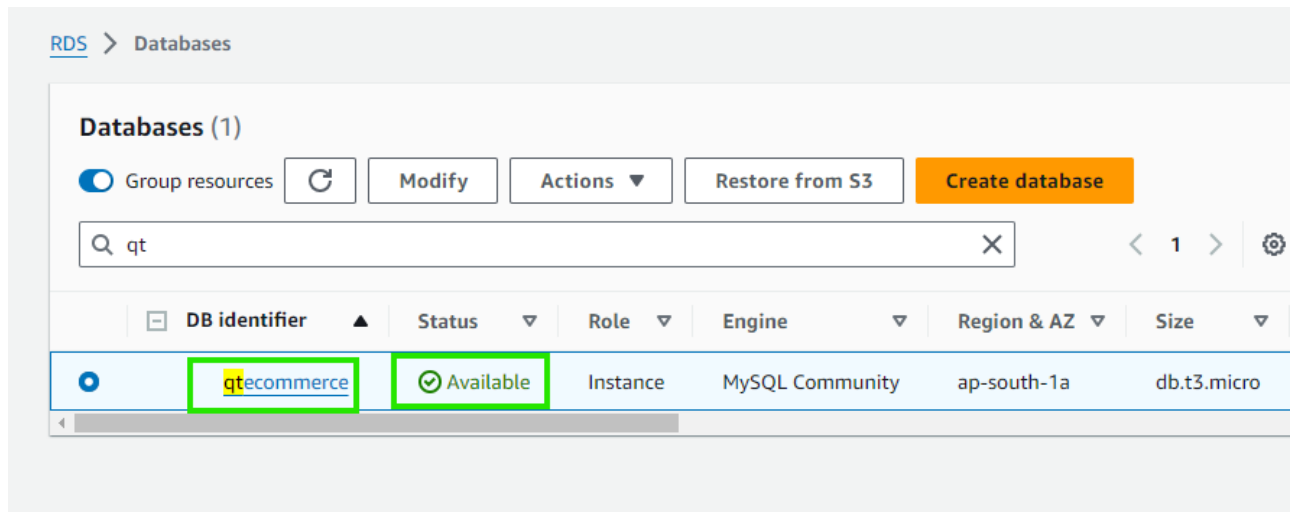
- Create a security group

A screenshot of the AWS VPC console. The left sidebar shows the 'VPC dashboard' with a search bar and a list of resources. The main area displays 'Security Groups (1/4) Info'. A table lists four security groups: 'ansible-insecure', 'rdsg', 'mysqlsg' (selected with a green box), and 'default'. The 'mysqlsg' group is associated with VPC ID 'vpc-06e49f59e4246fd88'. Below the table, the 'Inbound rules' tab is active, showing a single rule for 'MySQL/Aurora' on port 3306, also highlighted with a green box.

- Navigate to RDS



- For screenshots refer classroom video



- As discussed we can use mysql client or azure data studio

- [Refer Here](#) for creating a table and inserting data.

```
+-----+
| Tables_in_authors |
+-----+
| authors            |
+-----+
1 row in set (0.02 sec)

mysql> INSERT INTO authors (id,name,email) VALUES(1,"Vivek","xuz@abc.com");
Query OK, 1 row affected (0.02 sec)

mysql> INSERT INTO authors (id,name,email) VALUES(2,"Priya","p@gmail.com");
Query OK, 1 row affected (0.03 sec)


mysql> INSERT INTO authors (id,name,email) VALUES(3,"Tom","tom@yahoo.com");
Query OK, 1 row affected (0.02 sec)

mysql> SELECT * FROM authors;
+-----+-----+-----+
| id  | name | email      |
+-----+-----+-----+
| 1   | Vivek | xuz@abc.com |
| 2   | Priya | p@gmail.com |
| 3   | Tom   | tom@yahoo.com |
+-----+-----+-----+
3 rows in set (0.03 sec)


mysql>
```

## AWS Database Deployment Options


- Single Database



- Multi AZ Database



- Multi Cluster Database



## AWS Database Instance Sizes

- By Purpose
  - Standard
  - Memory Optimized
  - Compute Optimized

- Burstable
- Size [Refer Here](#)

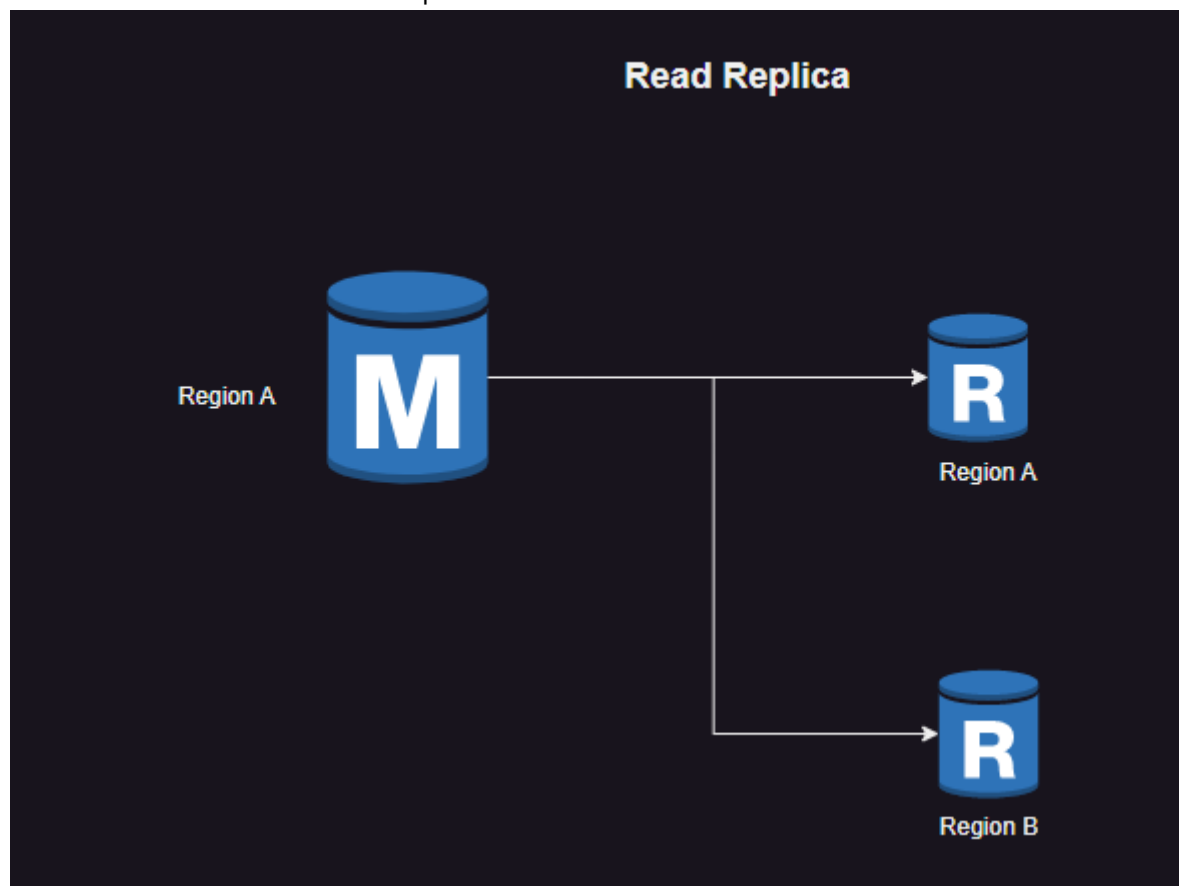
## Backups

- Manual Backups can be taken to create snapshots
- Automated Backups can be enabled.
  - PITR (Point in time restore)

## AWS RDS (Contd..)

### Read Replicas

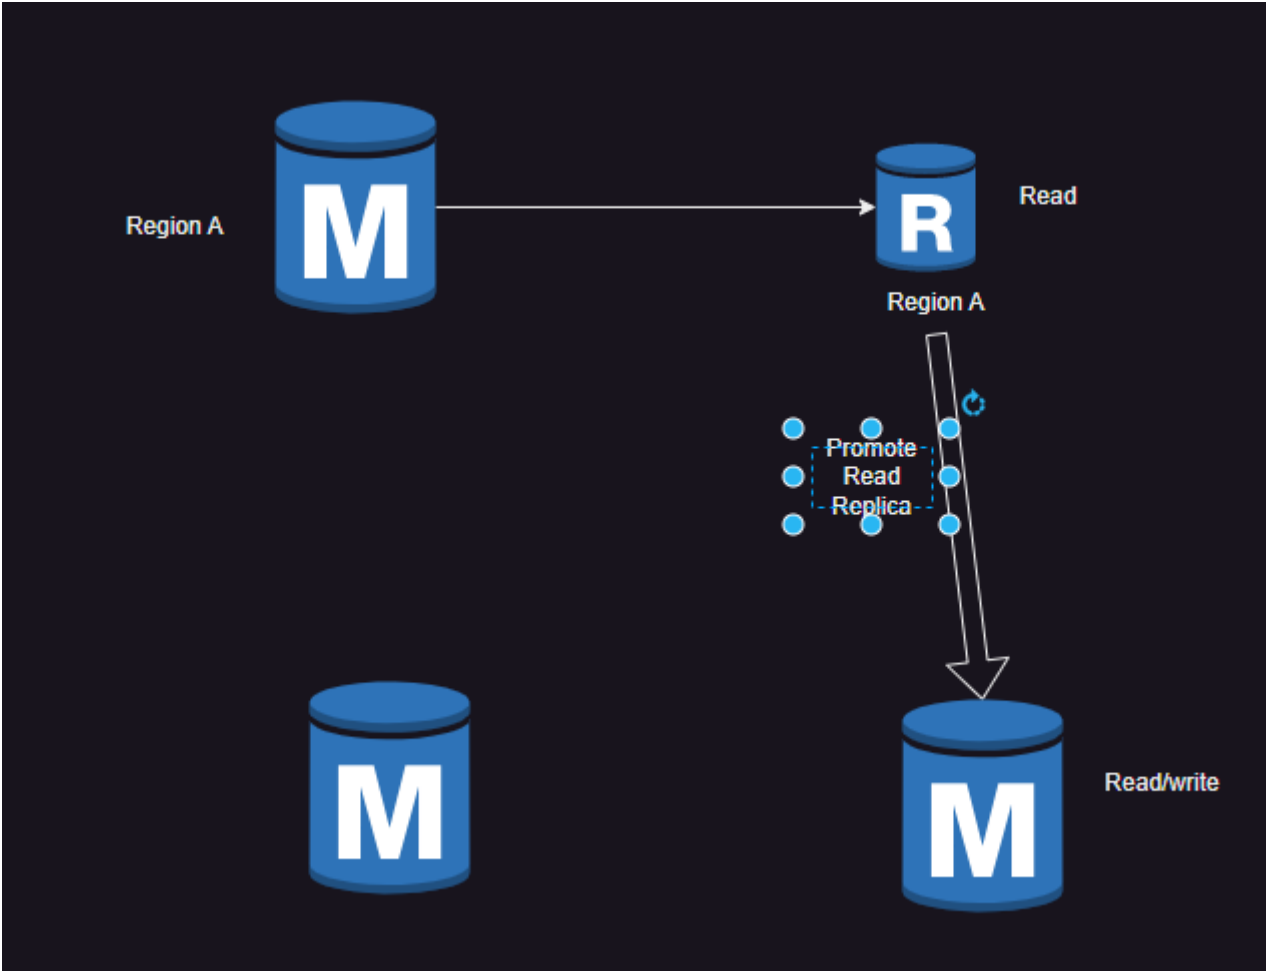
- AWS RDS allows to create read replicas



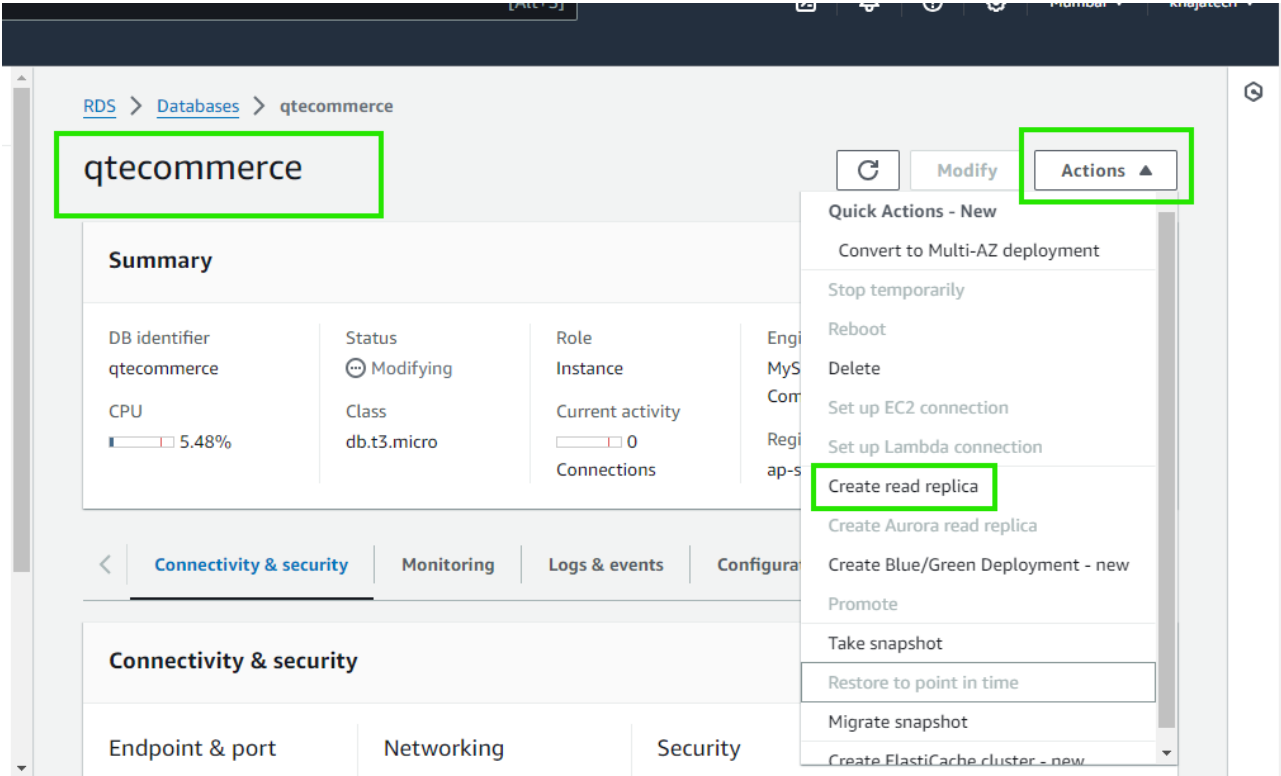
- Read Replicas can be in the same/different regions

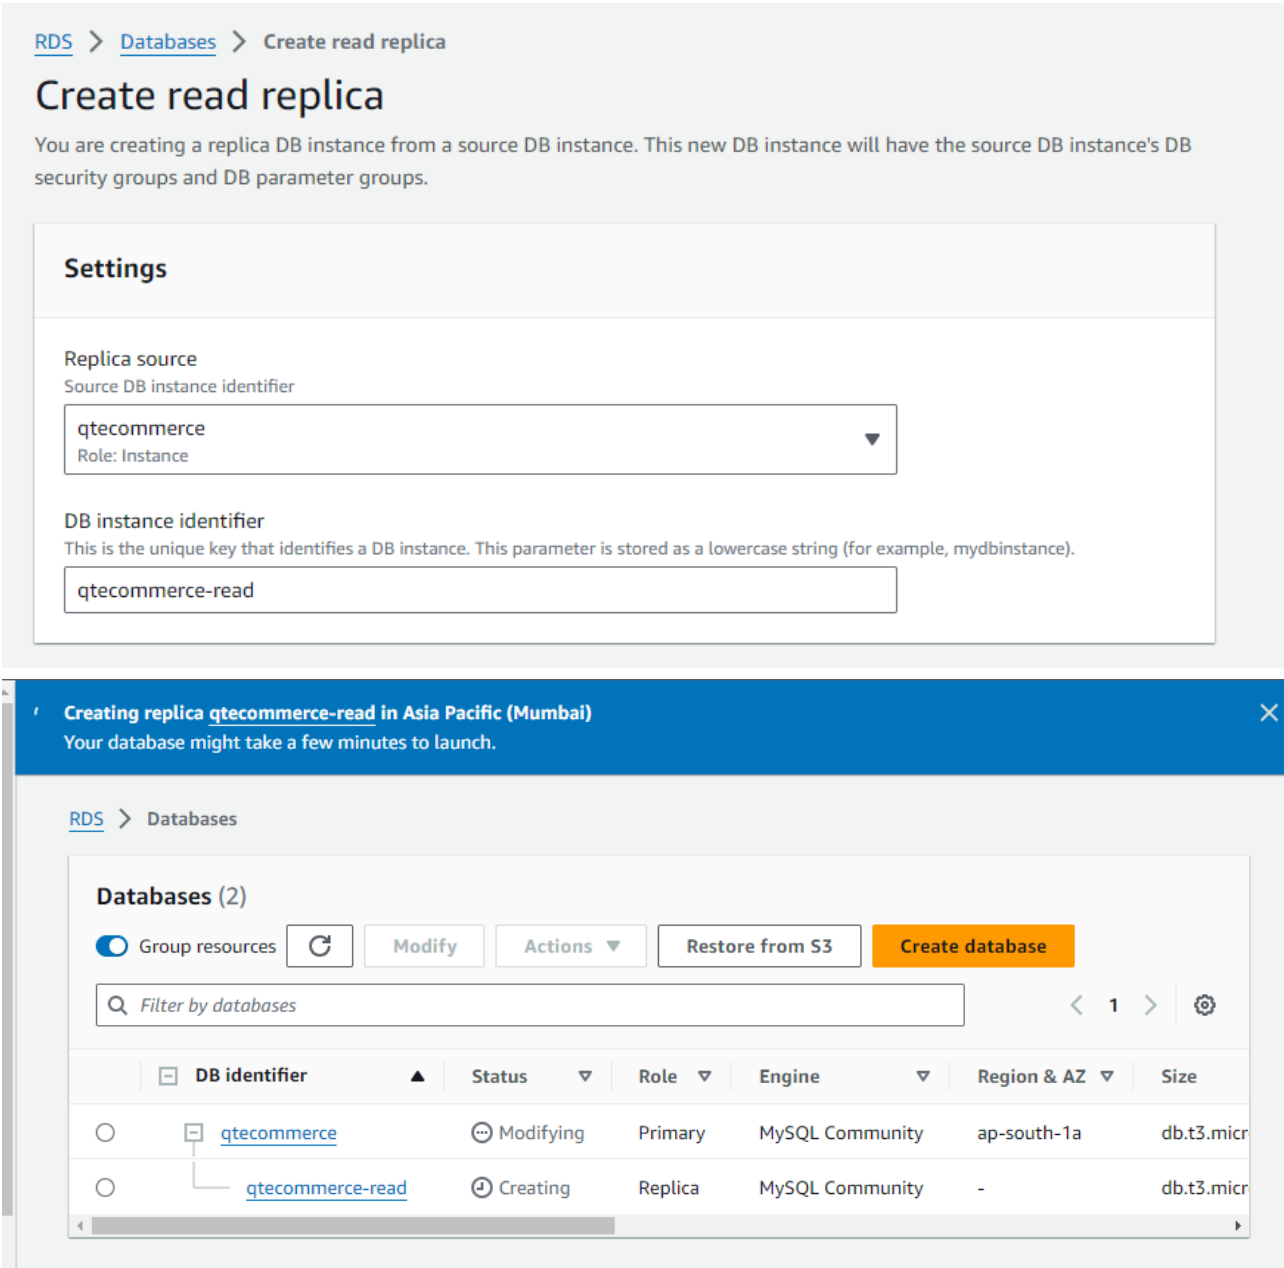


- Read Replicas can be promoted to accept write requests but the replication will no longer continue



- Ensure the Backups are on (atleast for 1 day)
- Create a Read Replica





- We have experimented with read replicas in the class watch video
- Multi Cluster DB Creation

Overview

Amazon RDS

Dashboard

Databases

Query Editor

Performance insights

Snapshots

Exports in Amazon S3

Automated backups

Reserved instances

Proxies

Subnet groups

Parameter groups

Option groups

Custom engine versions

Zero-ETL integrations

Successfully deleted DB instance gtecommerce-read

Notifications 0 0 2 0 2

RDS > Databases > database-1

database-1

Related

Filter by databases

DB identifier	Status	Role	Engine	Region & AZ
database-1	Creating	Multi-AZ DB cluster	MySQL Community	ap-south-1
database-1-instance-1	Creating	Writer instance	MySQL Community	ap-south-1a
database-1-instance-2	Creating	Reader instance	MySQL Community	ap-south-1c
database-1-instance-3	Creating	Reader instance	MySQL Community	ap-south-1b

Connectivity & security

Monitoring

Logs & events

Configuration

Maintenance & backups

Endpoints (2)

Find resources

Endpoint name	Status	Type	Port
database-1.cluster-c762uok2ezvw.ap-south-1.rds.amazonaws.com	Creating	Writer	3306
database-1.cluster-ro-c762uok2ezvw.ap-south-1.rds.amazonaws.com	Creating	Reader	3306

Connected compute resources (0)

Aurora

- Architecture  Preview

Exercise

- Create a mysql rds single database (db.t2.micro) with 2 days of backup
- Create a read replica
- check the connectivity
- delete both instances

## **Amazon Dynamo DB**

- To the Dynamo DB Table, we add indexes
- Indexes help in querying data
- There are two types of indexes
  - local secondary indexes (can be added during creation time only)
  - Global Secondary indexes (can be added post creation)

## **Amazon Document DB**

- Mongo DB as a service offered by AWS
- Document DB gives us options to run mongo much like RDS
- We have to deal with
  - security groups
  - subnet groups
  - vpcs
- Once created we can use mongoshell or application libraries to connect.

## **Other Databases**

- Cassandra ==> KeySpaces
- Graph => Neptune

## **Cache**

- Redis
  - OSS => Elastic Cache
  - Enterprise => Memory DB
- Memcached => Elastic Cache

## **Database Backups in AWS**

- Snapshot refers to a backup which is generally executed manually.
- AWS has a service to centralize all backups which is AWS Backup Service
- To Use AWS Backup Service, We need to
  - Create a vault
  - Backup plan
  - Protected Resources: Resources for which backup is enabled.

## **Backups in Azure**

-

[Azure](#) has Backup Center which indirectly creates Backup Vaults

Create a Backup for any Service

## **DataWarehouses**

- AWS Redshift: This is Datawarehouse as a service offered by AWS
- [Azure](#) Synapse: This is Datawarehouse as a service offered by Azure

## **Big Data**

- As we have lots of data in the non-queryable format, to use this big data technology has started
- Hadoop
- Spark
- Hive
- Cloud has started giving a storage called as blob storages (s3 and azure storage account).
- Organizations have started adopting blob storages as data lakes
- Organization have started building pipelines for data ingestion, storage and analytics
- Treating data as code and building data pipelines.